

TrustCheck verification dataset

3,750 labeled examples (1,250 per label). Task: given a claim and an evidence passage, predict the verdict.

Fields

Field	Meaning
id	Unique row id
category	capitals, elements_number, elements_symbol, events, books, planets, geography, formulas, science
claim	The statement to verify
evidence	The retrieved passage the verdict must be based on
label	supported, uncertain, unsupported
reasoning	Short explanation, usable as a generation target

Field	Meaning
split	train (about 80%), val (about 9%), test (about 10%)

Splits are assigned by claim subject, so the same subject never appears in two splits.

Labeling rules (match prd.md and aistack.md)

- supported: the evidence clearly states what the claim says.
- unsupported: the evidence states a different value, so it contradicts the claim.
- uncertain: the evidence is about something else (insufficient), or two sources conflict.

Limits

- Synthetic and templated from a fixed fact table of about 300 well-known facts. Good for bootstrapping, format testing and a baseline classifier. It will not teach a model real-world evidence reading on its own.
- Evidence passages are written, not scraped from real sources.

- Before trusting results, add real claim and evidence pairs from your own TrustCheck runs and hold them out as the test set.

Regenerate or change size: edit PER_LABEL in `build_dataset.py` and run `python3 build_dataset.py` .